

Governance Documentation

This document defines the governance controls for the AI agent used in the automation process. Its purpose is to outline authorised actions and rigid guardrails that the AI agent must follow, in line with risks and non-negotiables discussed with the stakeholder. The aim of this document is to define and maintain legal accuracy, transparency and traceability of the solution.

Agent mandate

The agent is authorised to:

- Retrieve newly published judgments from the Find Case Law API.
- Read judgment metadata (e.g. title, court, publication date, tags).
- Read the judgment text for analysis.
- Assess the judgment against the configured monitoring topics.
- Store actions to do with processing judgments, along with associated decisions in the database.

The agent operates strictly as a decision-support tool, not an autonomous decision-maker. All outputs produced by the agent must remain subject to human review before any further use, internal or external.

The agent is not authorised to:

- Make the final decision on whether a judgment is relevant (approval/rejection).
- Publish or distribute case summaries.
- Provide legal advice or legal opinions.
- Modify, delete or overwrite source judgments.
- Change monitoring topics, prompts or system configuration.
- Access systems or data outside those explicitly required for the workflow.
- Execute actions that are not explicitly defined within the workflow.

Audit logging

The solution maintains a comprehensive audit trail of every significant action performed throughout the lifecycle of a judgment. Audit records are stored in an AWS RDS database, to where both automated actions and human decisions are logged.

Audit entries are logged at explicit points in the process, when:

- a judgment is ingested
- a judgment is filtered
- relevance is assigned
- a human reviewer approves or rejects a relevance assessment at triage review
- a briefing is generated from a judgment
- a human reviewer approves or rejects a briefing at briefing review

Each audit entry captures the information required to reconstruct the processing history of a judgment. This includes:

- the timestamp of the event
- the case name
- the event
- the actor responsible (system or user)
- the decision or action performed
- a reason or justification where appropriate.

This information is displayed through the History tab on the dashboard. Users can view a chronological timeline of all recorded events through a searchable interface. The history page allows users to filter records by the above fields, making it straightforward to investigate the lifecycle of a particular judgment or review organisational activity over a specified period.

Human-in-the-loop gates

The agent has two human-in-the-loop gates built into its process: these are points at which a human (in this case, a Case Desk member) signs off on the agent's outputs and decisions up to that point. These gates are blocking events, meaning the agent cannot continue with the process until the human decision (approval/rejection) has been made.

The first human-in-the-loop gate is after case triage. After a new judgment has been identified and assessed, a Case Desk member reviews the proposed relevance before deciding whether the judgment should progress for summarisation or not. This ensures that the decision to process a judgment in its entirety and generate a briefing falls with a human reviewer, allowing contextual judgement and Case Desk experience to be applied where the agent's assessment may not fully capture the significance or audience value of a case.

The second human-in-the-loop gate takes place after briefing material has been generated. A Case Desk member reviews the summary to verify that it accurately reflects the judgment and meets JusticeLensUK's standards. This approval stage provides quality

assurance before the briefing is passed on for senior review, and ensures that responsibility for published content remains with the Case Desk team rather than the automated system.